

Qwen2.5-7B-Instruct | Quantized layer 14 | float16 $\rightarrow$ int8