

Qwen2.5-7B-Instruct | Quantized layer 24 | float16 → int4

