

Qwen2.5-7B-Instruct | Quantized layer 23 | float32 → int16

