

Qwen2.5-7B-Instruct | Quantized layer 24 | float32  $\rightarrow$  int16