

Qwen2.5-7B-Instruct | Quantized layer 26 | float32 → int16

1e-7+1

