

Qwen2.5-7B-Instruct | Quantized layer 25 | float32  $\rightarrow$  int16

