

Qwen2.5-7B-Instruct | Quantized layer 4 | float32 → int8

