

Qwen2.5-7B-Instruct | Quantized layer 16 | float32 $\rightarrow$ int8