

Qwen2.5-7B-Instruct | Quantized layer 8 | float32 → int4

