

Qwen2.5-7B-Instruct | Quantized layer 0 | float16  $\rightarrow$  int4

