

Qwen2.5-7B-Instruct | Quantized layer 16 | float32 → int4

