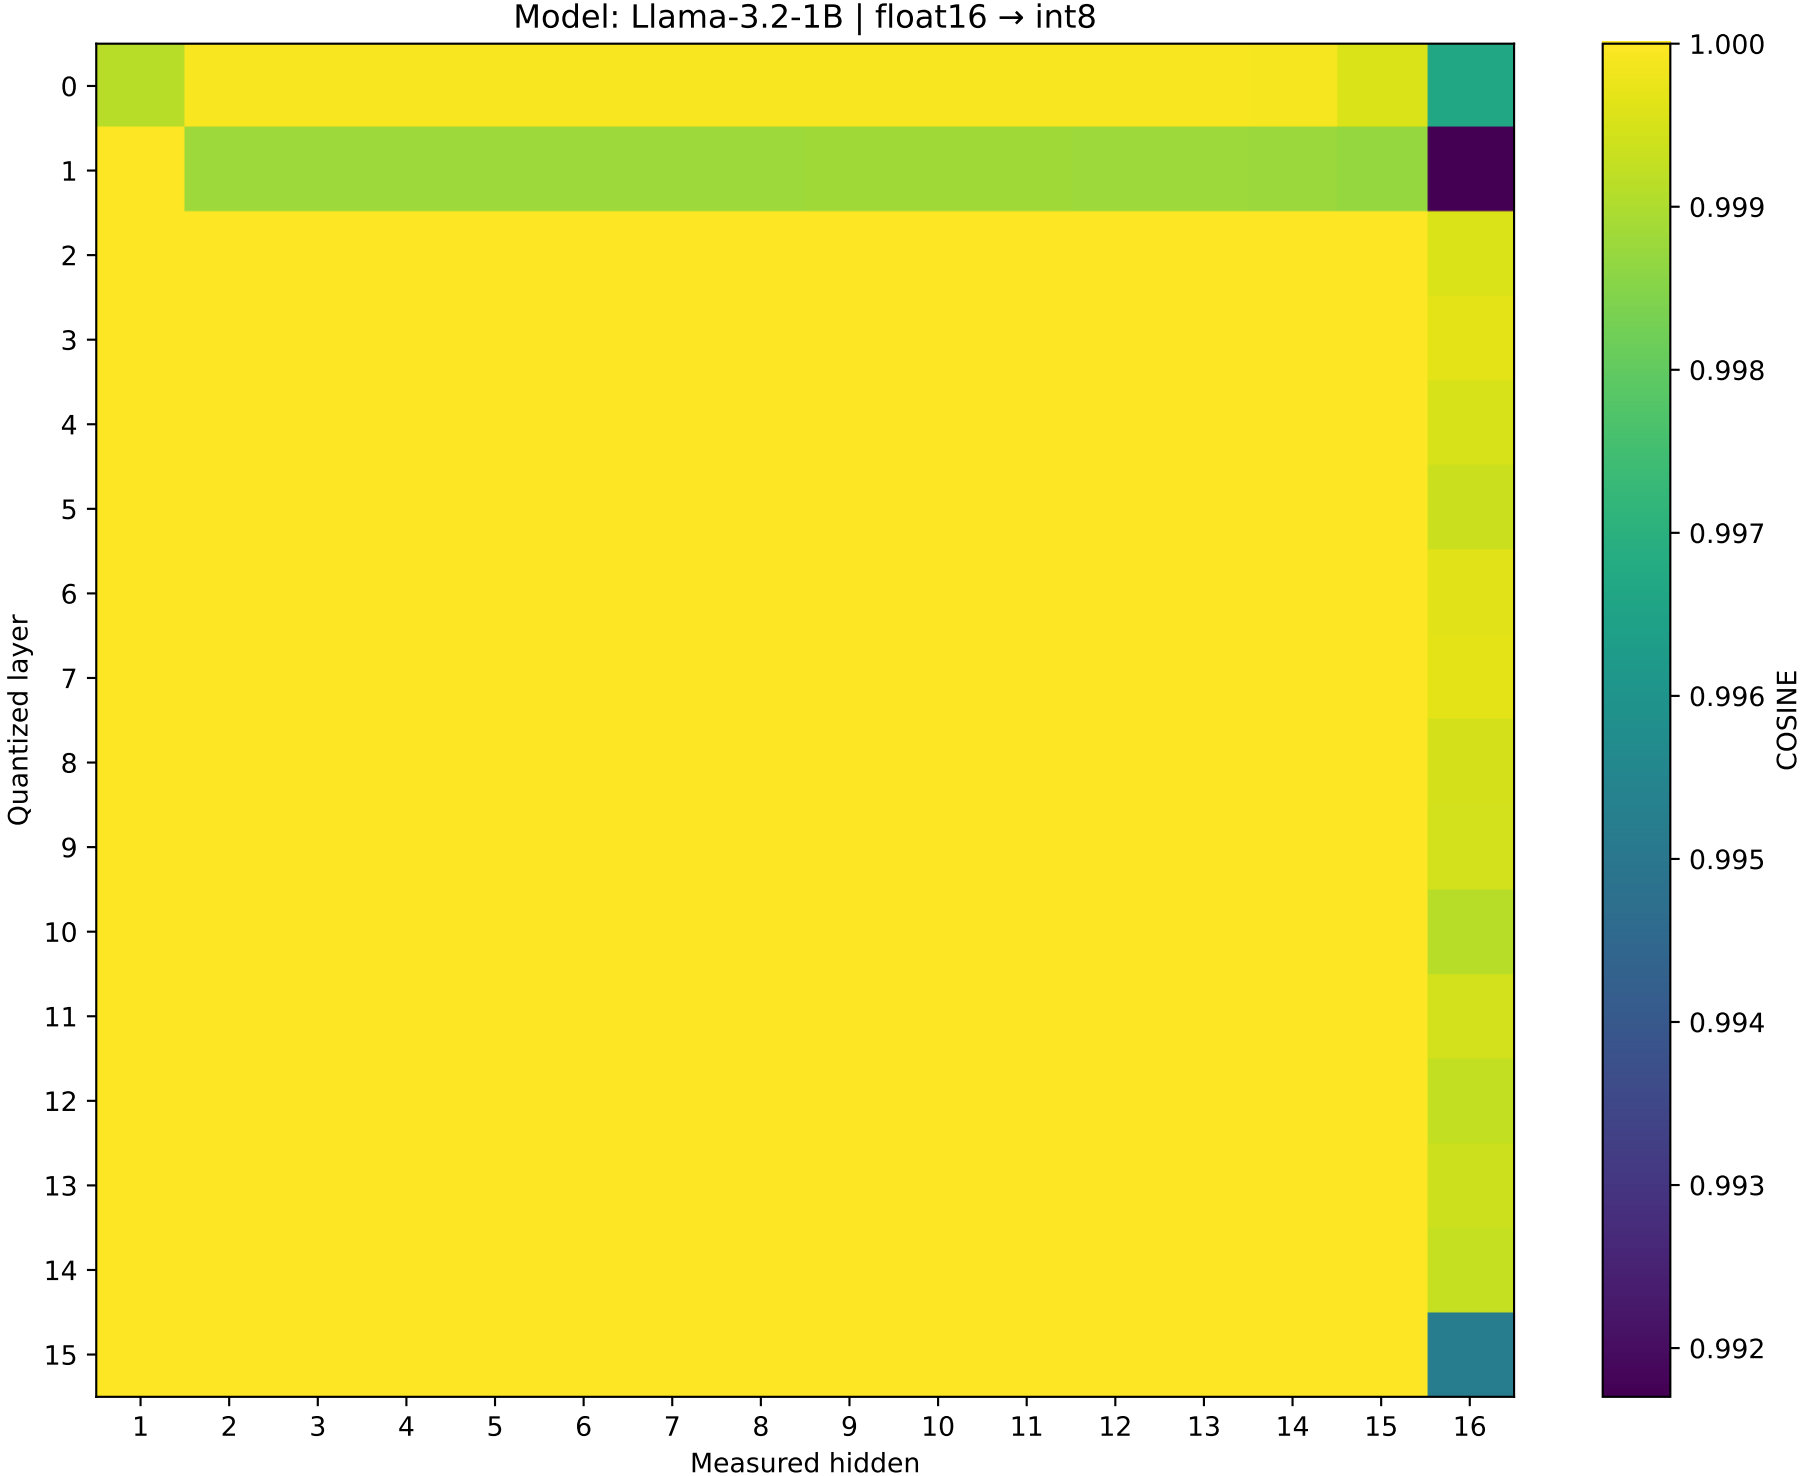

Model: Llama-3.2-1B | float16 → int8
Quantized layer
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
Measured hidden
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
COSINE
1.000
0.999
0.998
0.997
0.996
0.995
0.994
0.993
0.992