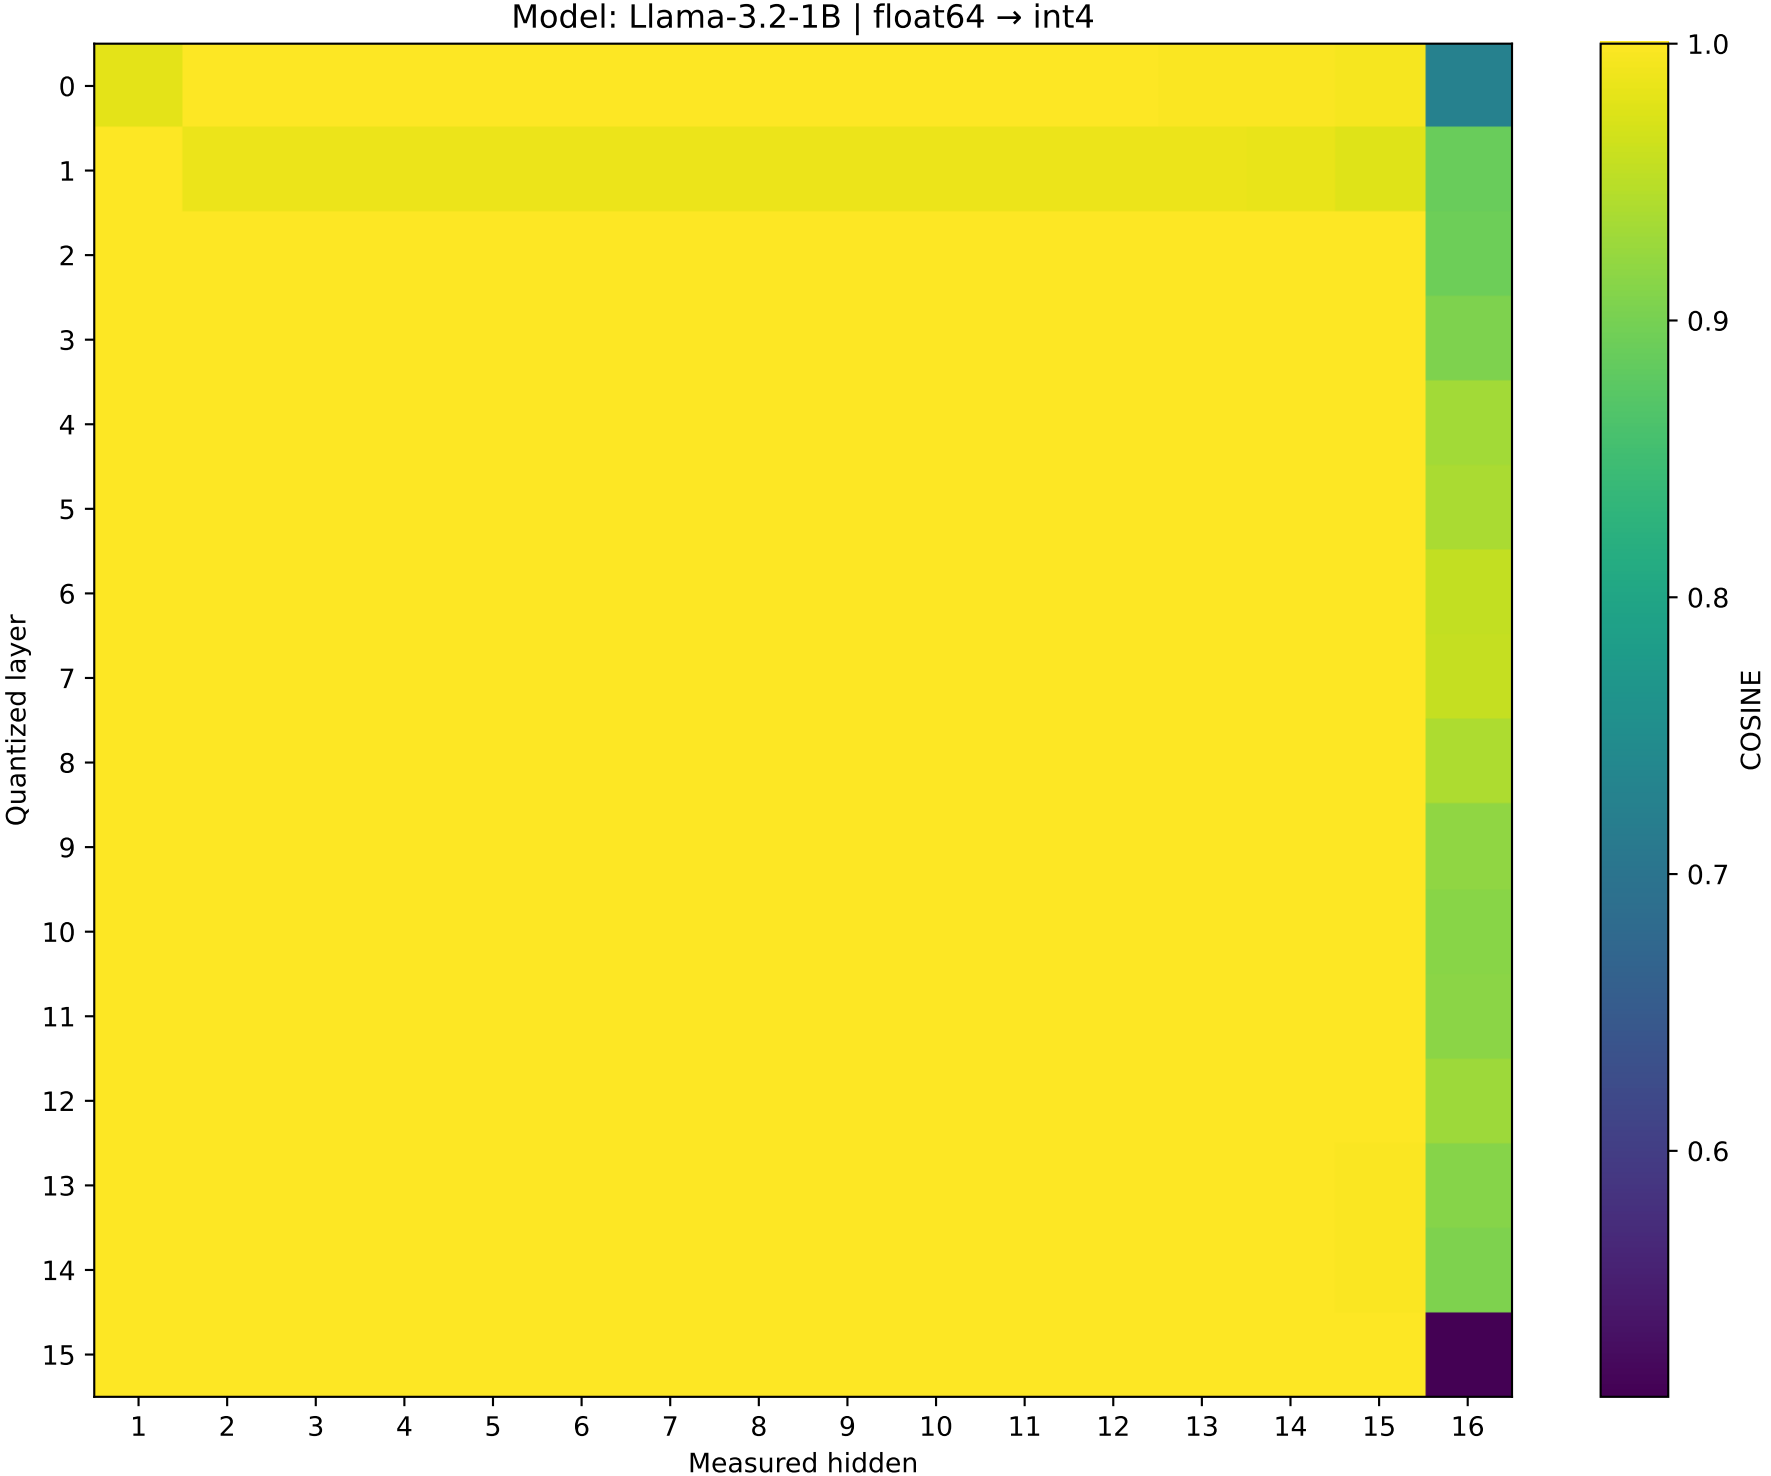

Model: Llama-3.2-1B | float64 → int4
Quantized layer
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
Measured hidden
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
COSINE
1.0
0.9
0.8
0.7
0.6