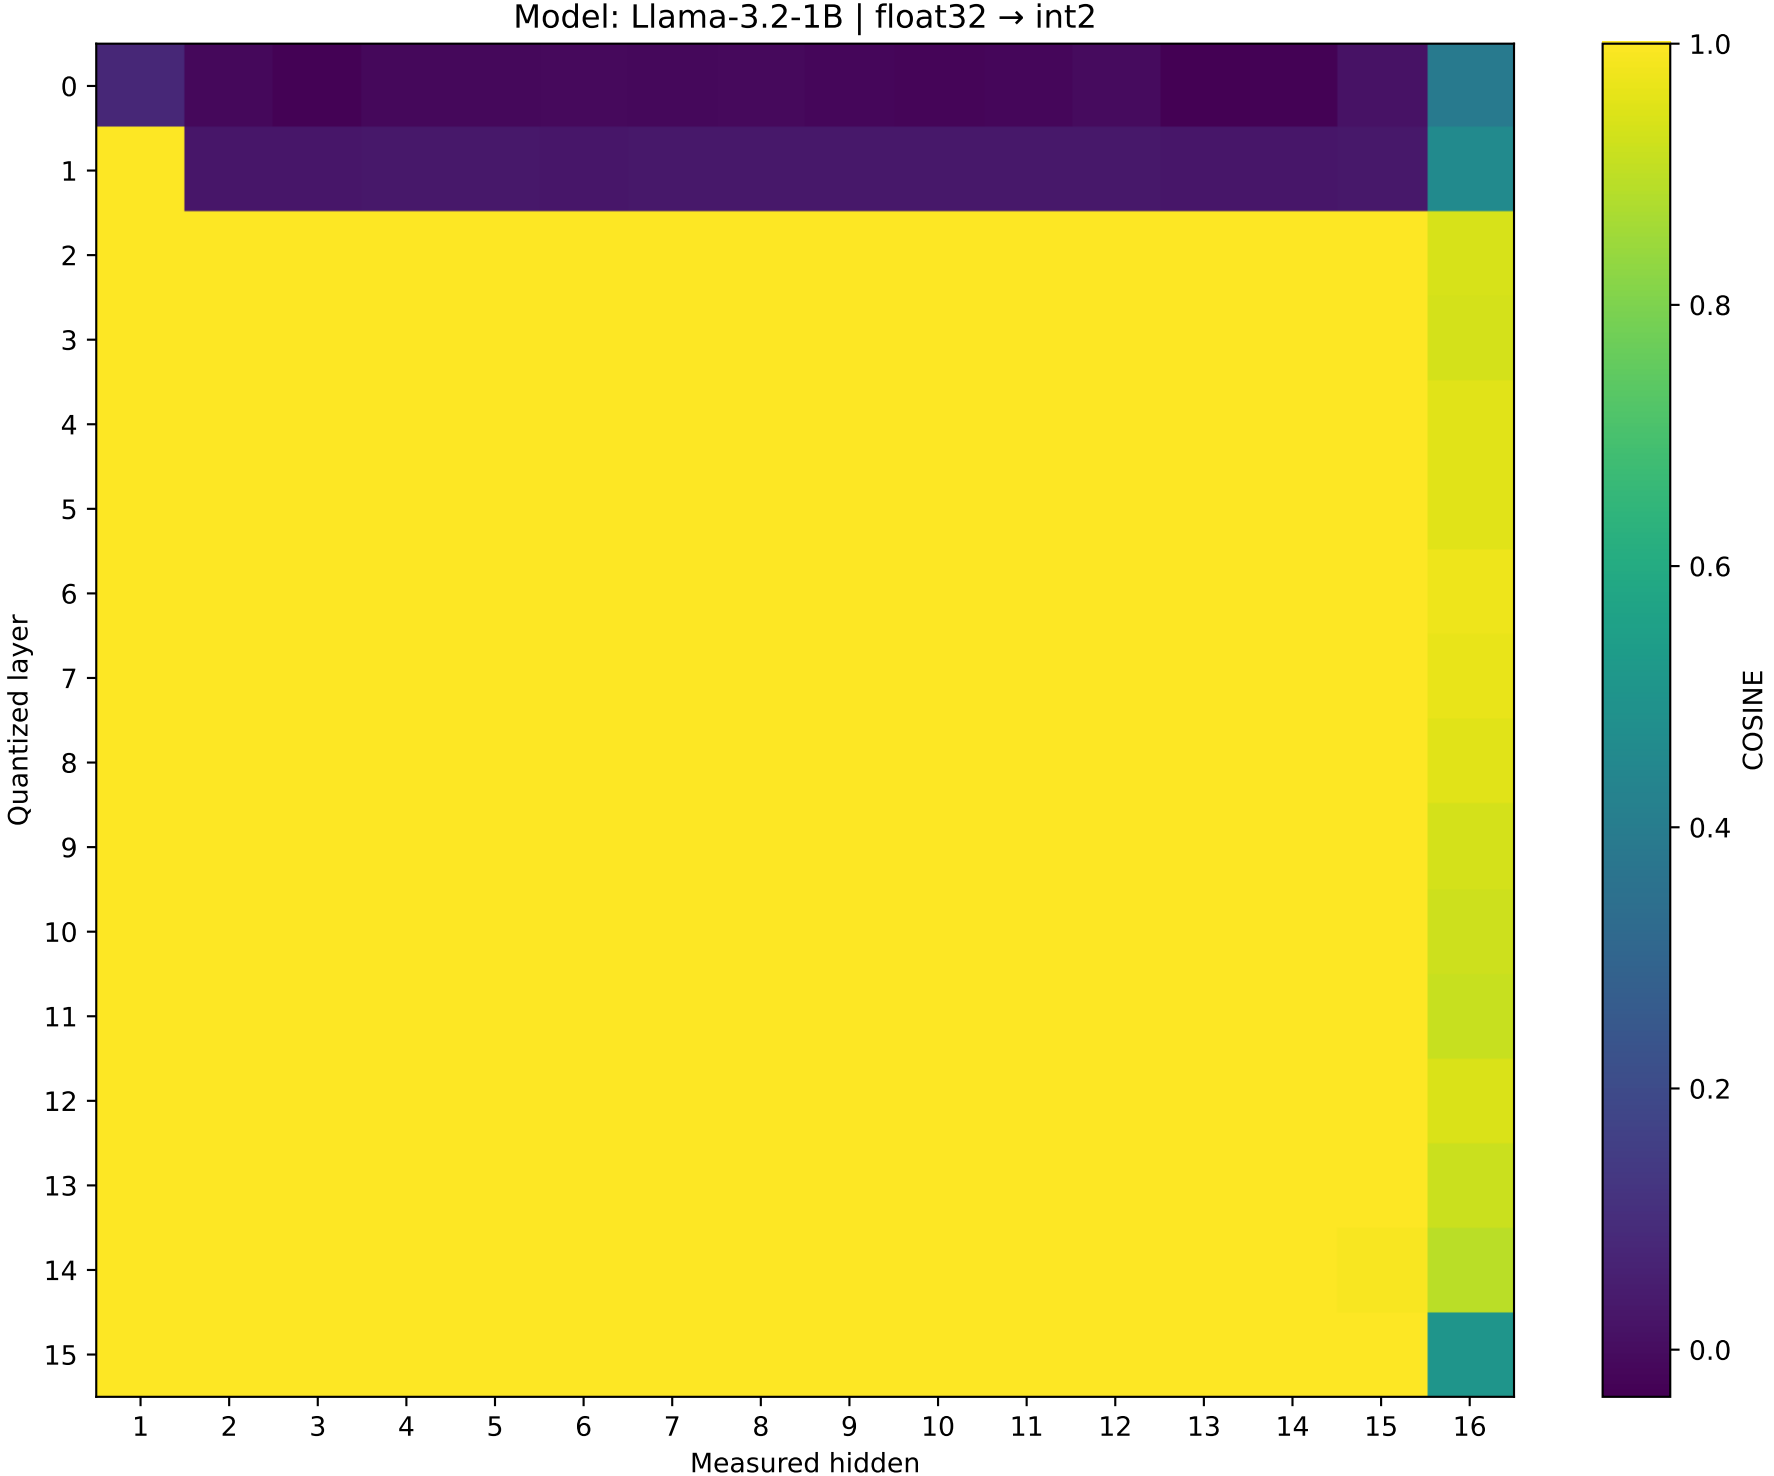

Model: Llama-3.2-1B | float32 → int2
Quantized layer
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
Measured hidden
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
COSINE
1.0
0.8
0.6
0.4
0.2
0.0