

Llama-3.2-3B | Quantized layer 24 | float32 → int16

