

Llama-3.2-3B | Quantized layer 8 | float32 → int16

